

# Parallelization in Python —Dask

Arnav Sood

NYU

April 21, 2016

# Dask Hierarchy

**Data Structures:** Jobs fit easily into operations on common data structures.

# Dask Hierarchy

**Data Structures:** Jobs fit easily into operations on common data structures.

**Imperatives:** Job is more complex, but doesn't warrant an entire custom graph.

# Dask Hierarchy

**Data Structures:** Jobs fit easily into operations on common data structures.

**Imperatives:** Job is more complex, but doesn't warrant an entire custom graph.

**Custom Graphs:** Intricate job.

# Dask Hierarchy

**Data Structures:** Jobs fit easily into operations on common data structures.

**Imperatives:** Job is more complex, but doesn't warrant an entire custom graph.

**Custom Graphs:** Intricate job.

**Schedulers:** Take the above specs, and run as efficiently as possible.

# Chunked Algorithms

```
# Try chunked matrix operations
import dask.array as da
import numpy as np
baseArray = np.random.rand(10000,10000)
# Make positive-definite
baseArray = np.dot(baseArray,baseArray.transpose())
smallChunks = da.from_array(baseArray,chunks=(100))
largeChunks = da.from_array(baseArray,chunks=(1000))
```

# Chunked Algorithms

```
# Try chunked matrix operations
import dask.array as da
import numpy as np
baseArray = np.random.rand(10000,10000)
# Make positive-definite
baseArray = np.dot(baseArray,baseArray.transpose())
smallChunks = da.from_array(baseArray,chunks=(100))
largeChunks = da.from_array(baseArray,chunks=(1000))
```

```
In [2]: %time np.linalg.cholesky(baseArray)
        %time da.linalg.cholesky(smallChunks)
        %time da.linalg.cholesky(largeChunks)
```

```
CPU times: user 21.5 s, sys: 849 ms, total: 22.4 s
```

```
Wall time: 6.43 s
```

```
CPU times: user 385 ms, sys: 26.4 ms, total: 412 ms
```

```
Wall time: 440 ms
```

```
CPU times: user 444  $\mu$ s, sys: 1  $\mu$ s, total: 445  $\mu$ s
```

```
Wall time: 449  $\mu$ s
```

# Lazy Evaluation

```
np0nes = np.ones((10000,10000))  
%time np.exp(np0nes)[1:100,1:100]
```

```
dask0nes = da.ones((10000,10000), chunks=(100))  
%time da.exp(dask0nes)[1:100,1:100]
```



# Lazy Evaluation

```
npOnes = np.ones((10000,10000))  
%time np.exp(npOnes)[1:100,1:100]
```

```
daskOnes = da.ones((10000,10000), chunks=(100))  
%time da.exp(daskOnes)[1:100,1:100]
```

**Result:** 10× speedup. (970ms → 90ms)

# Ghosting

# Internals: Dictionary

```
import itertools slice =  
itertools.islice(smallChunks.dask.items(), 0, 1) for  
key, value in slice: print(key,value)
```

# Internals: Dictionary

```
import itertools slice =  
itertools.islice(smallChunks.dask.items(), 0, 1) for  
key, value in slice: print(key,value)
```

## Result:

**Key:** from-array-8c463a3b962efd2e81b47e1832acfb81', 83, 94)

**Value:** (<function getarray at 0x1083260d0>,  
'from-array-8c463a3b962efd2e81b47e1832acfb81', (slice(8300,  
8400, None), slice(9400, 9500, None)))

# API

**Subset** of NumPy API.

# API

**Subset** of NumPy API.  
Good things:

# API

**Subset** of NumPy API.

Good things:

Bad things:

# API

**Subset** of NumPy API.

Good things:

Bad things:

Reference: `dir(dask.array)`



# API

**Subset** of NumPy API.

Good things:

Bad things:

Reference: `dir(dask.array)`